


Docker and CoreOS

A bit of what why and how


Hi I'm Jim!



docker



Core OS

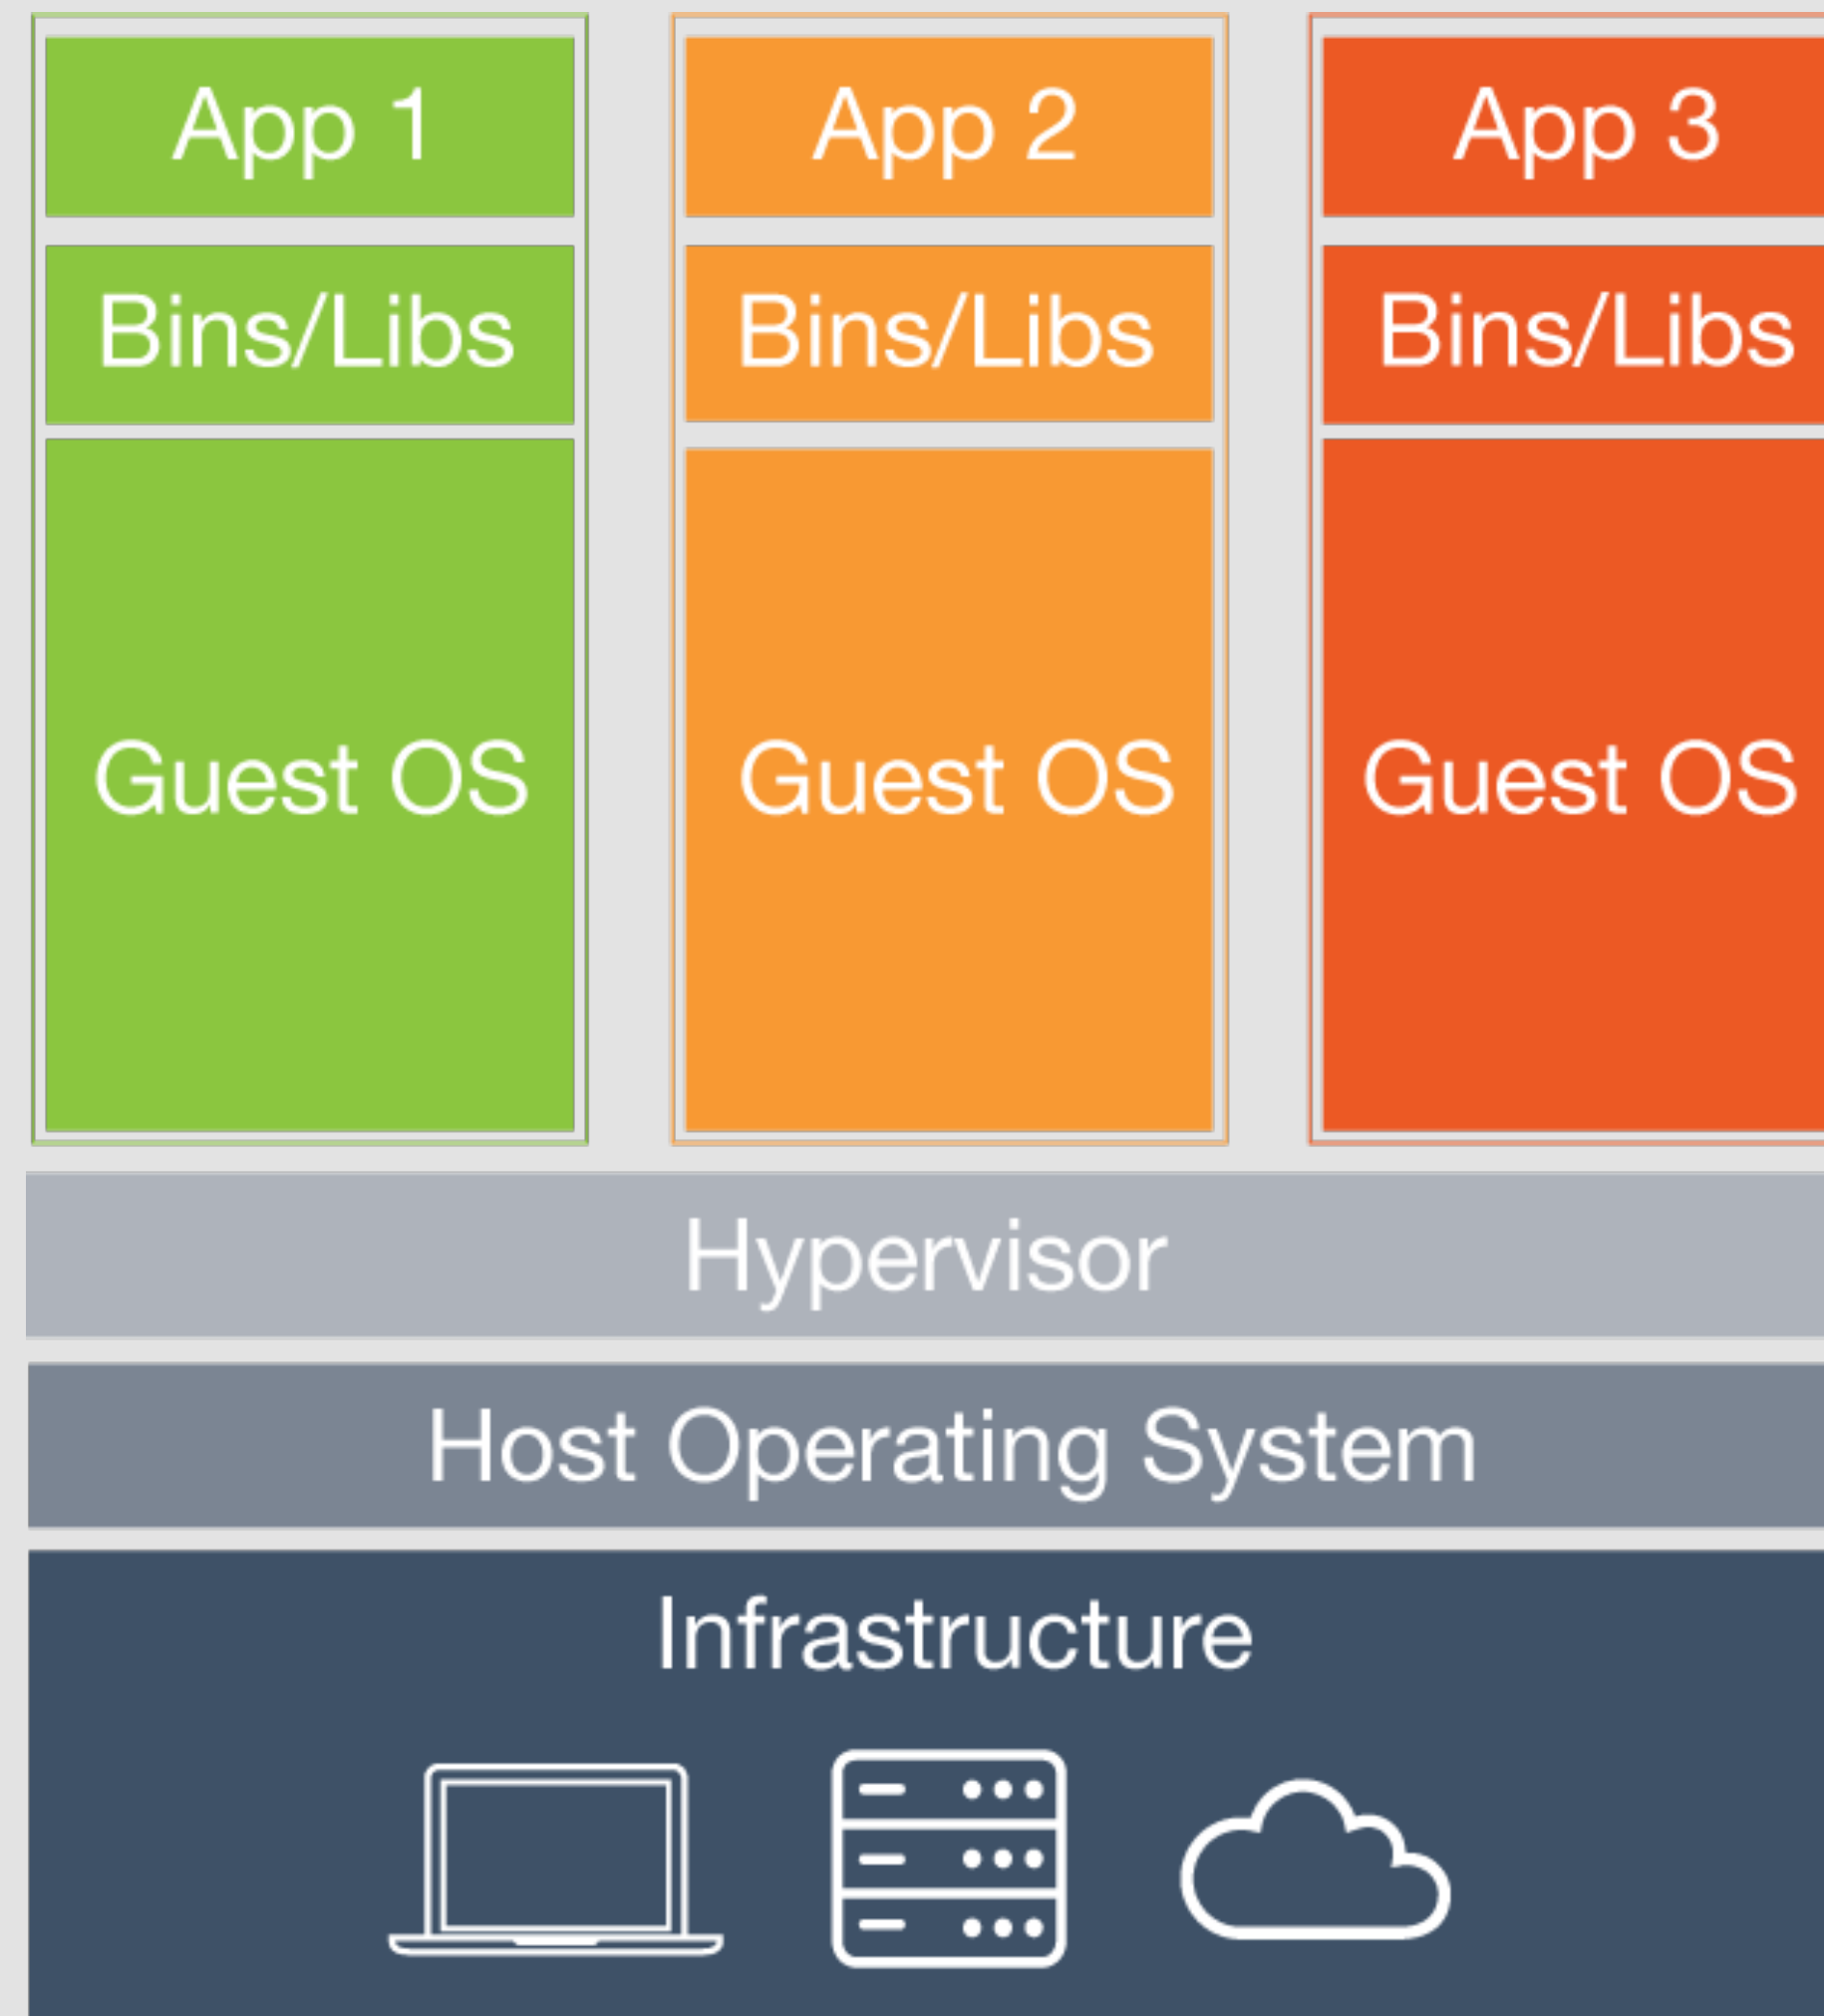


docker

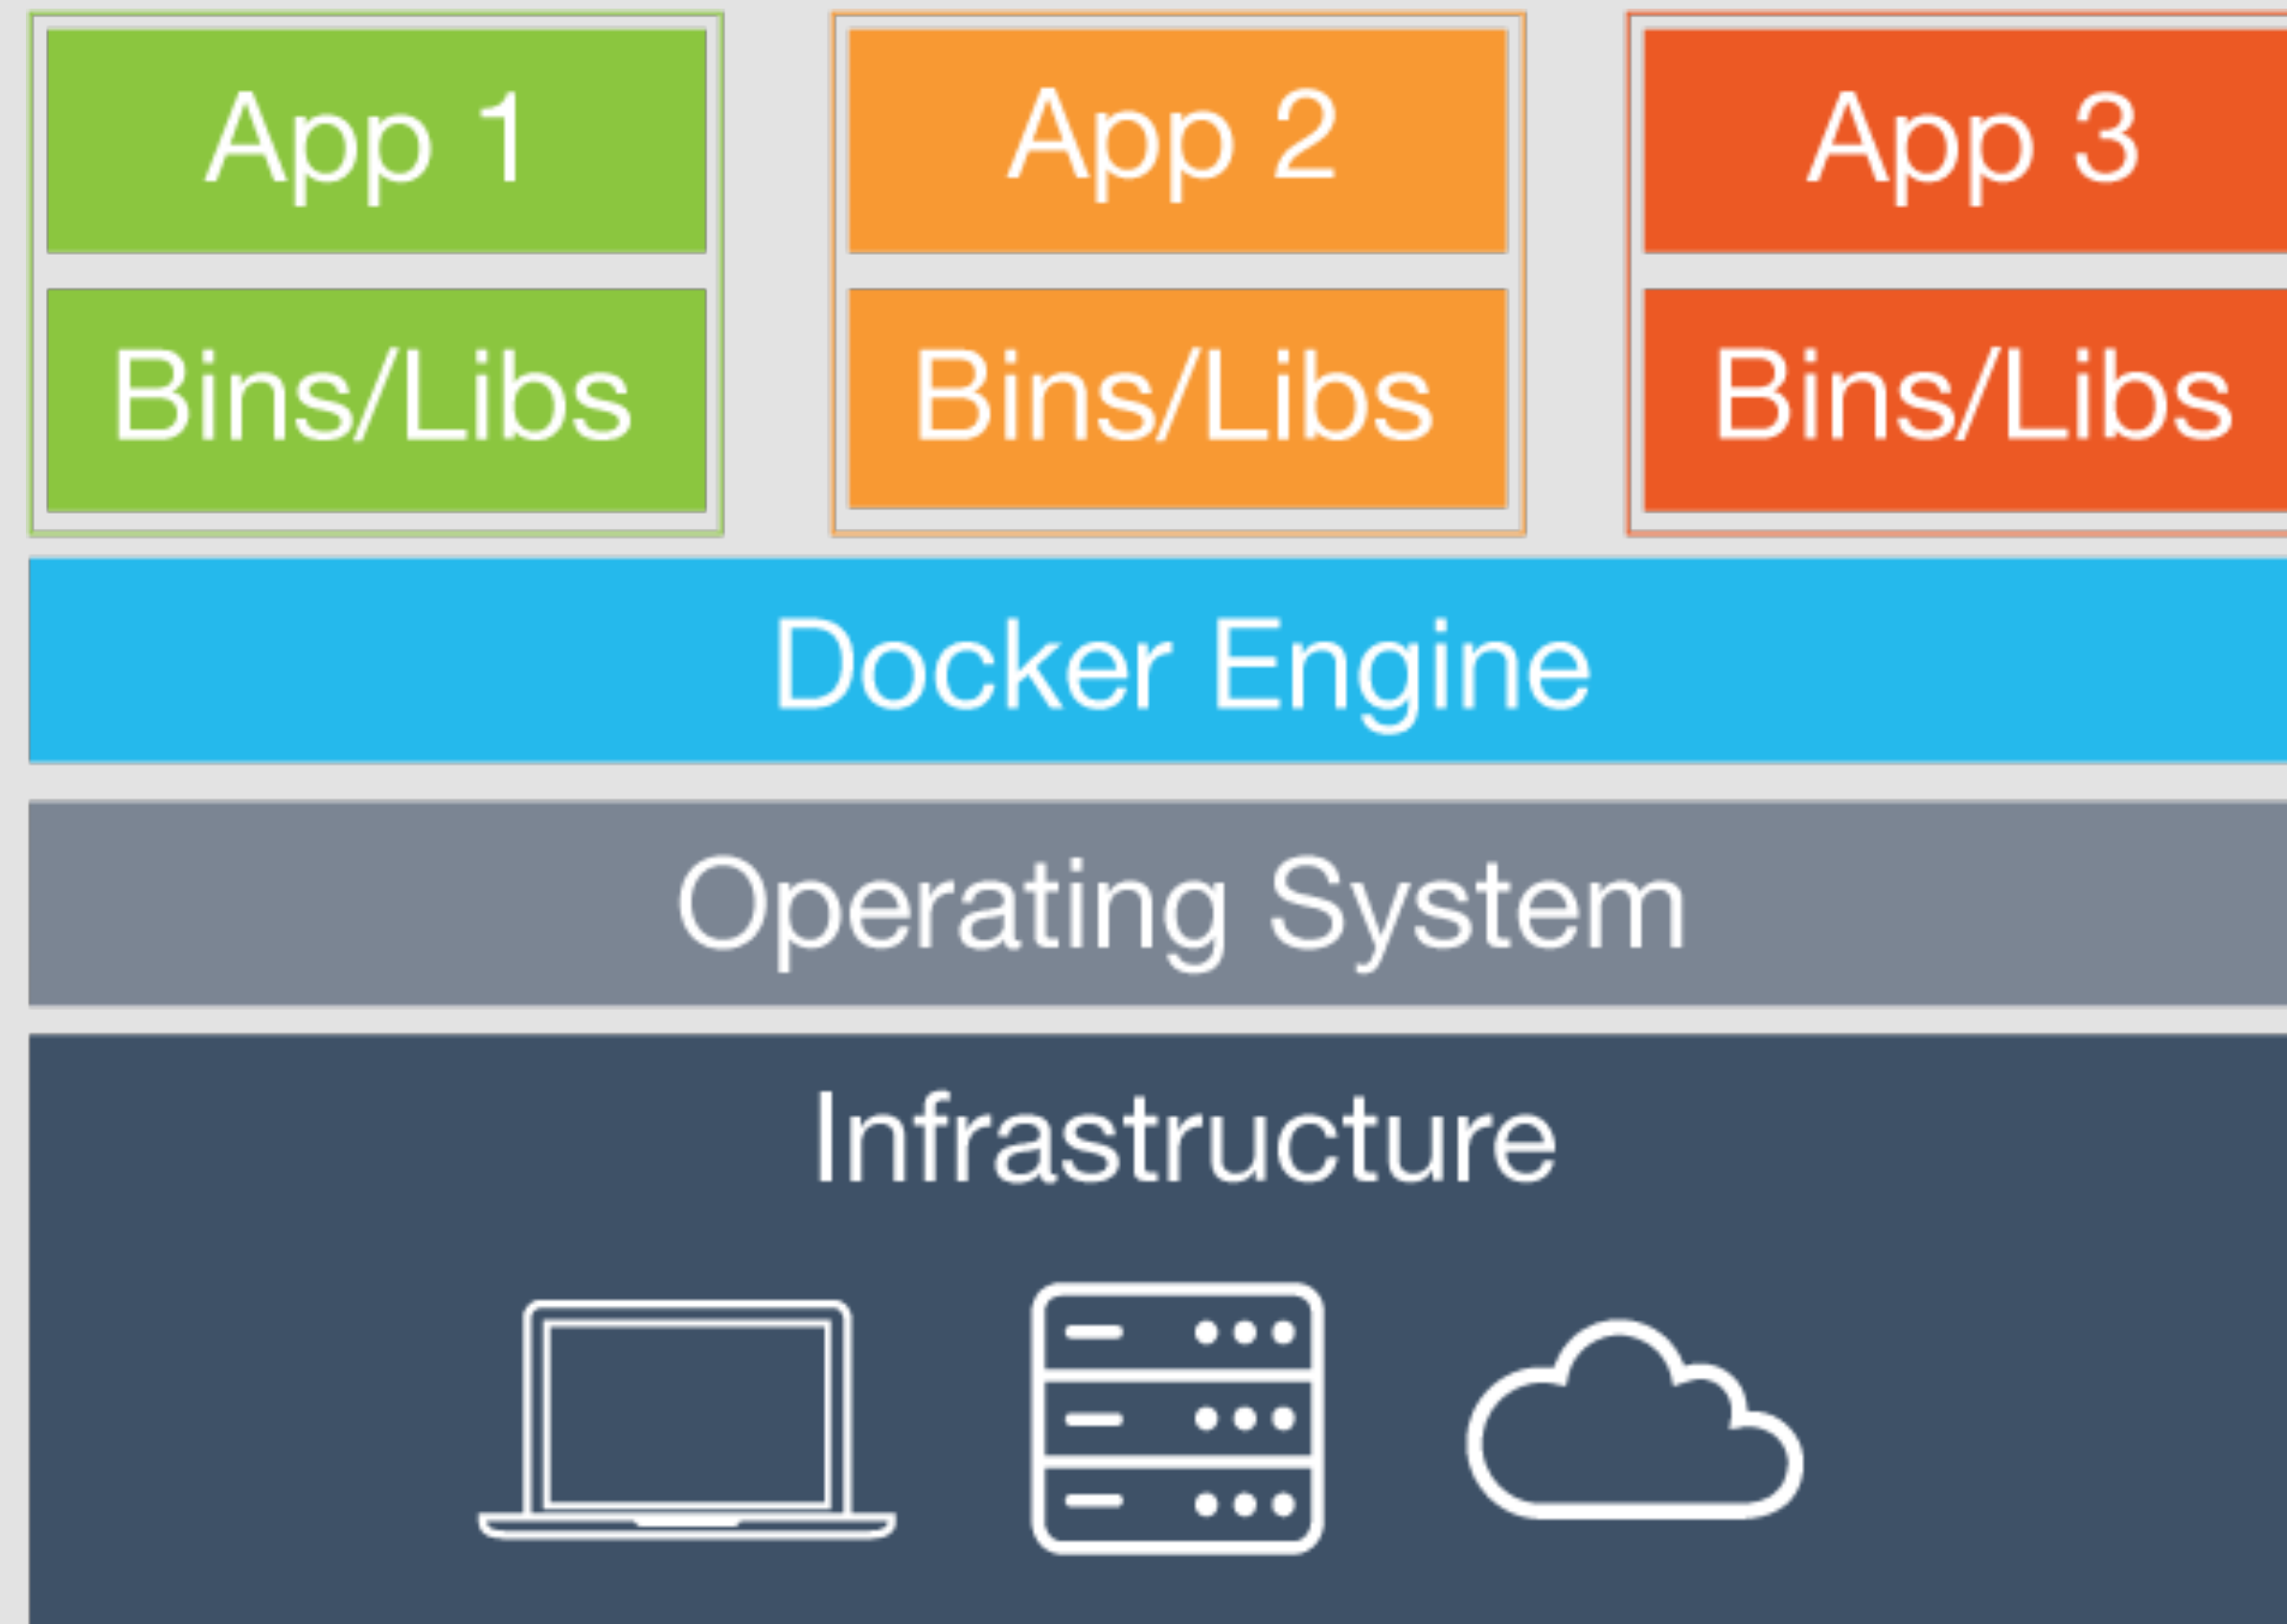
**YO DAWG WE HEARD YOU LIKED DOCKER
SO WE PUT A DOCKER IN YOUR DOCKER**

**SO YOU CAN DOCKER
WHILE YOU DOCKER**

VMs




Containers



Some Docker Terminology

- A **Dockerfile** is a text document that contains all the commands to assemble an image
- An **Image** is a shareable snapshot of software
- A **Container** is a running instance of an image
- A **Registry** is a repo of docker images

How these parts relate



Dockerfile

```
FROM golang:latest  
MAINTAINER Jim Weber
```

```
RUN mkdir -p /go/src/valet/conf
```

```
COPY *.go /go/src/valet/  
COPY ./conf/* /go/src/valet/conf/
```

```
RUN go get github.com/pkg/profile  
RUN cd /go/src/valet ; go build
```

```
CMD /go/src/valet/valet -P 8000
```

```
EXPOSE 8000
```

Dockerfile

```
FROM golang:latest  
MAINTAINER Jim Weber
```

```
RUN mkdir -p /go/src/valet/conf
```

```
COPY *.go /go/src/valet/  
COPY ./conf/* /go/src/valet/conf/
```

```
RUN go get github.com/pkg/profile  
RUN cd /go/src/valet ; go build
```

```
CMD /go/src/valet/valet -P 8000
```

```
EXPOSE 8000
```

Dockerfile

```
FROM golang:latest  
MAINTAINER Jim Weber
```

```
RUN mkdir -p /go/src/valet/conf
```

```
COPY *.go /go/src/valet/  
COPY ./conf/* /go/src/valet/conf/
```

```
RUN go get github.com/pkg/profile  
RUN cd /go/src/valet ; go build
```

```
CMD /go/src/valet/valet -P 8000
```

```
EXPOSE 8000
```

Dockerfile

```
FROM golang:latest  
MAINTAINER Jim Weber
```

```
RUN mkdir -p /go/src/valet/conf
```

```
COPY *.go /go/src/valet/  
COPY ./conf/* /go/src/valet/conf/
```

```
RUN go get github.com/pkg/profile  
RUN cd /go/src/valet ; go build
```

```
CMD /go/src/valet/valet -P 8000
```

```
EXPOSE 8000
```

Dockerfile

```
FROM golang:latest  
MAINTAINER Jim Weber
```

```
RUN mkdir -p /go/src/valet/conf
```

```
COPY *.go /go/src/valet/  
COPY ./conf/* /go/src/valet/conf/
```

```
RUN go get github.com/pkg/profile  
RUN cd /go/src/valet ; go build
```

```
CMD /go/src/valet/valet -P 8000
```

```
EXPOSE 8000
```

Dockerfile

```
FROM golang:latest  
MAINTAINER Jim Weber
```

```
RUN mkdir -p /go/src/valet/conf
```

```
COPY *.go /go/src/valet/  
COPY ./conf/* /go/src/valet/conf/
```

```
RUN go get github.com/pkg/profile  
RUN cd /go/src/valet ; go build
```

```
CMD /go/src/valet/valet -P 8000
```

```
EXPOSE 8000
```

Let's build an image

**Ok ... But how do
I use this in practice?**

NGINX

**Ruby 1.9
App**

**Ruby 2.1
App**

**Java 6
App**

**Java 9
App**

**Node 3
App**

**Node 5
App**

**Go
App**

**Swift
App**

I'm a Server running docker I don't need

- compilers
- language run times
- special Libraries

NGINX

Ruby 1.9
App

Ruby 2.1
App

Java 6
App

Java 9
App

Node 3
App

Node 5
App

Go
App

Swift
App

**Do one thing
and, do it well**

**Do one thing
and, do it well**

It is the UNIX philosophy too



**But Jim, How do I
use docker in an
HA Cluster?**

**YO DAWG, I HEARD YOU LIKE
VIRTUALIZATION**

**SO I RUN DOCKER IN A VM WHICH RUNS A
VIRTUAL MACHINE TO RUN CONTAINERS**

memegenerator.net

Ok, not really



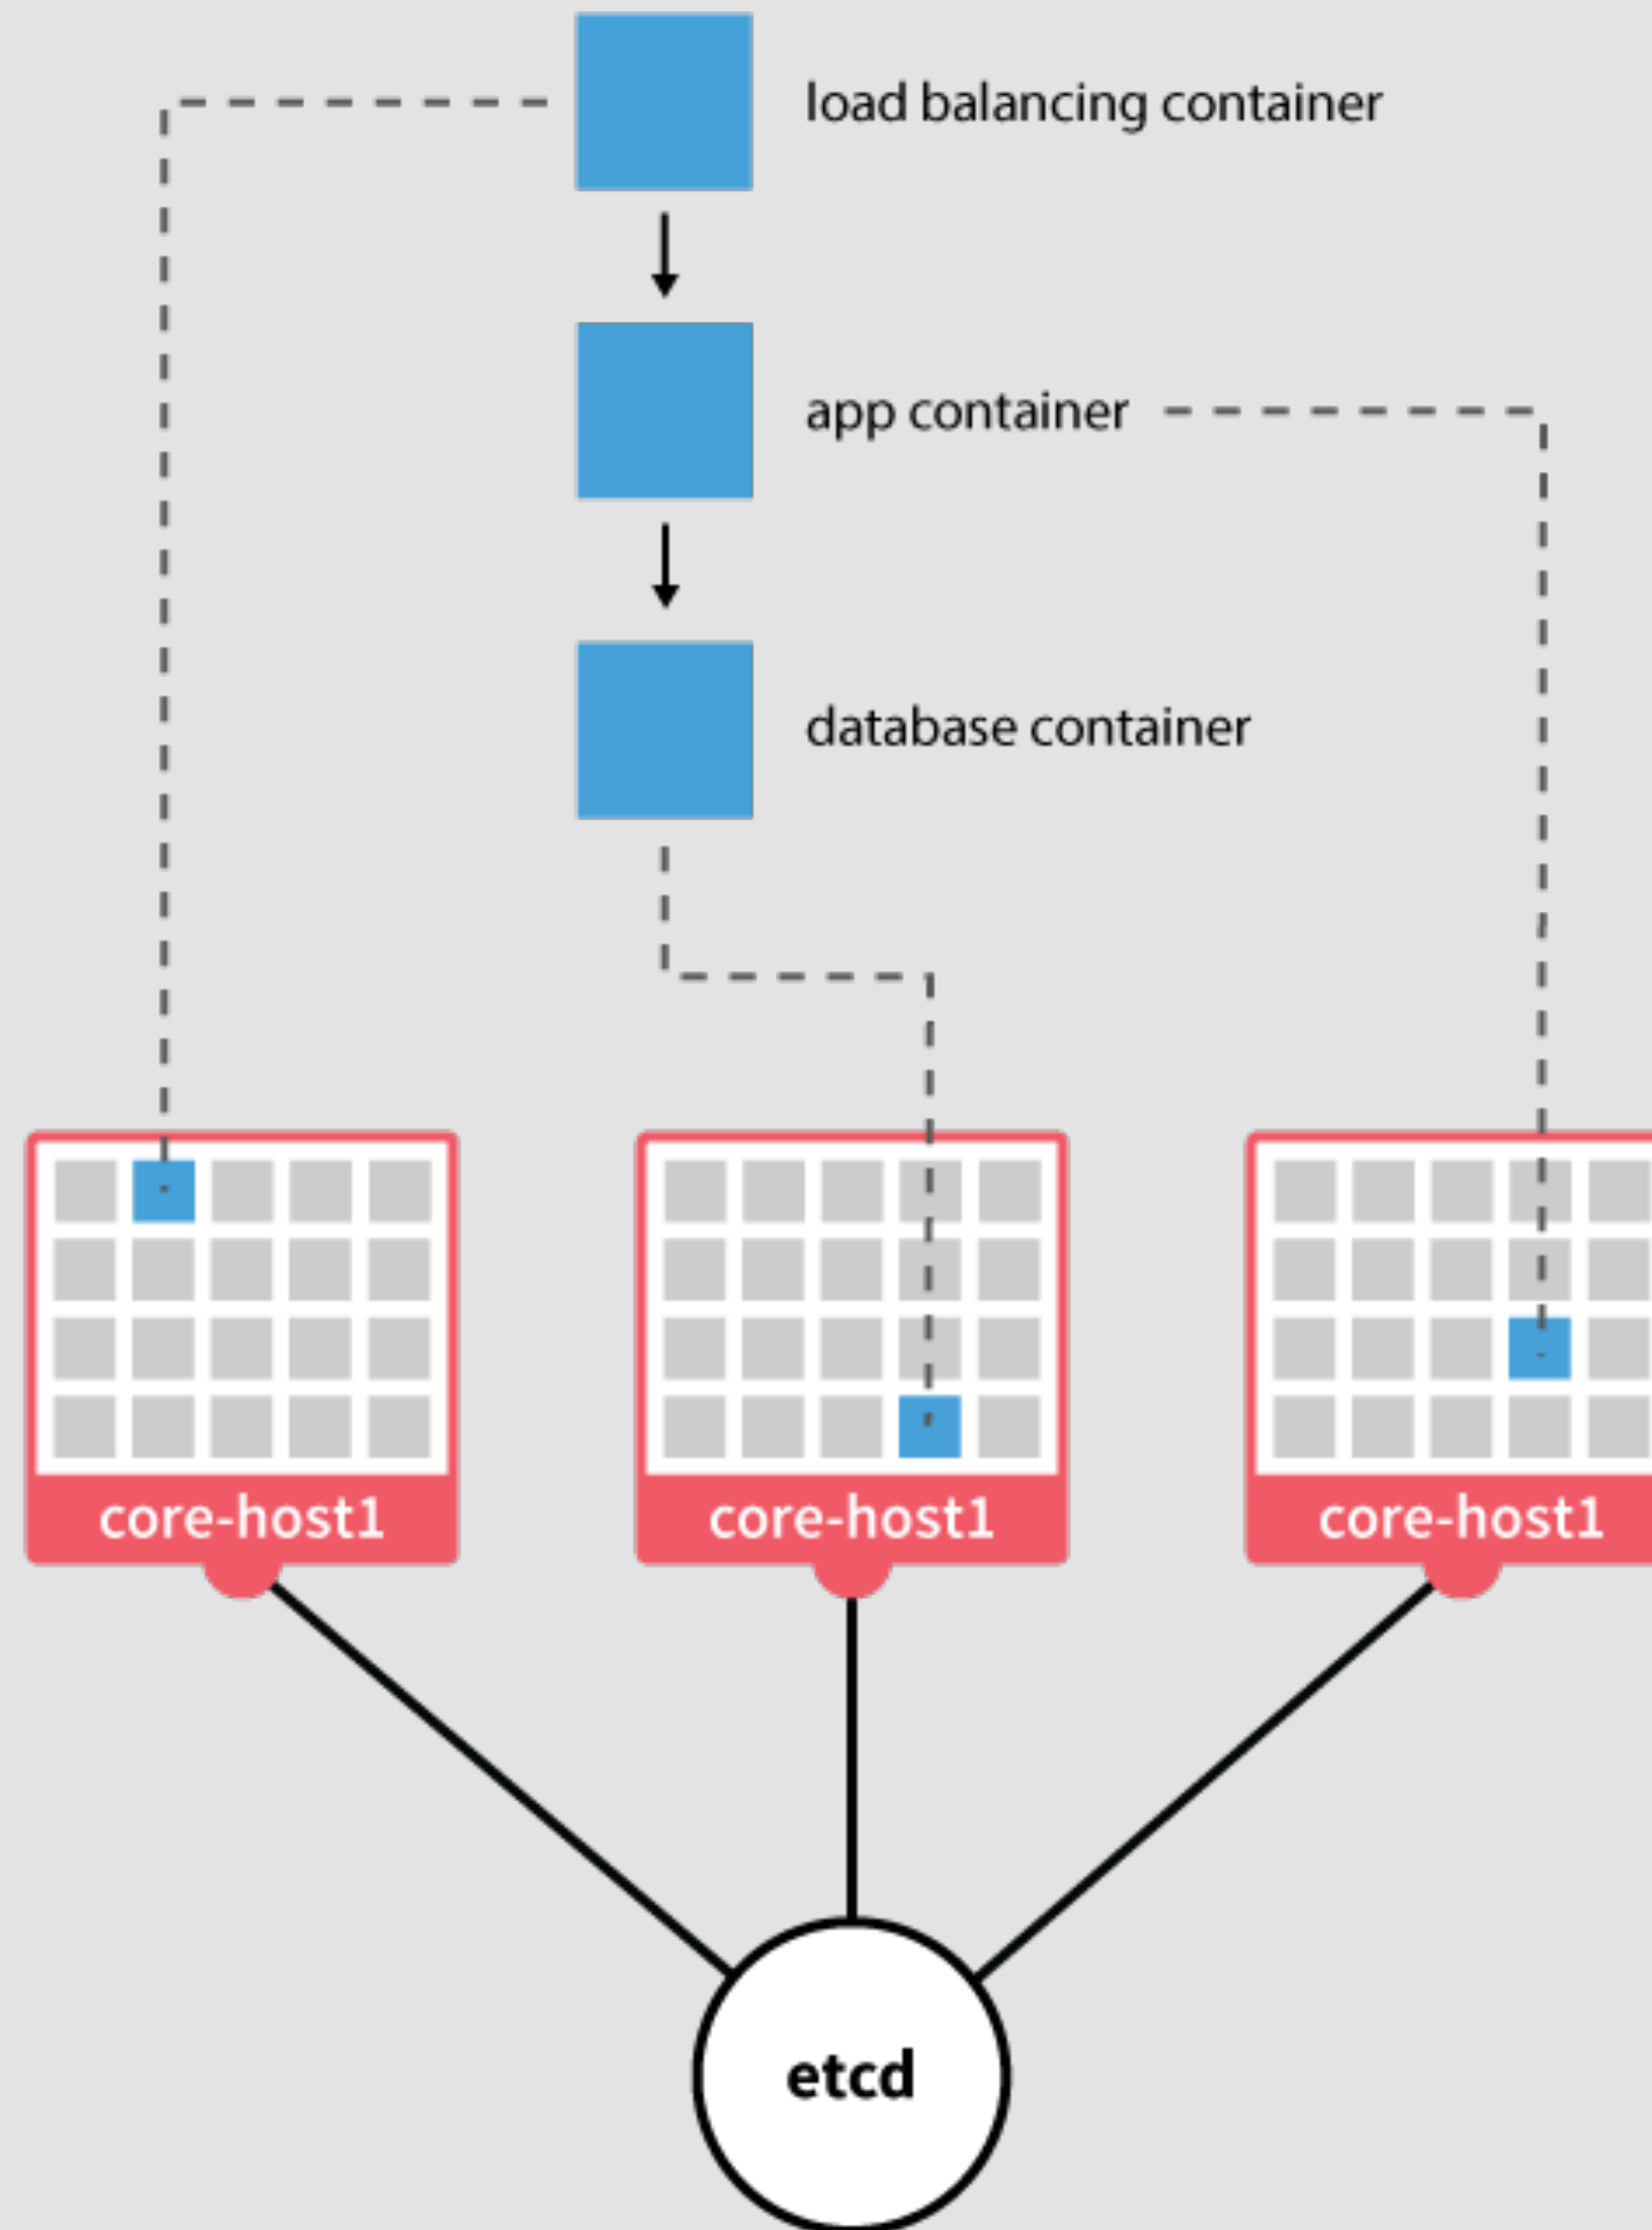
Core OS




Core OS

CoreOS is a powerful Linux distribution built to make large, scalable deployments on varied infrastructure simple to manage

A Linux distro made specifically for running
containers in a clustered environment.
With some extra tools to make that work.



confd



etcd

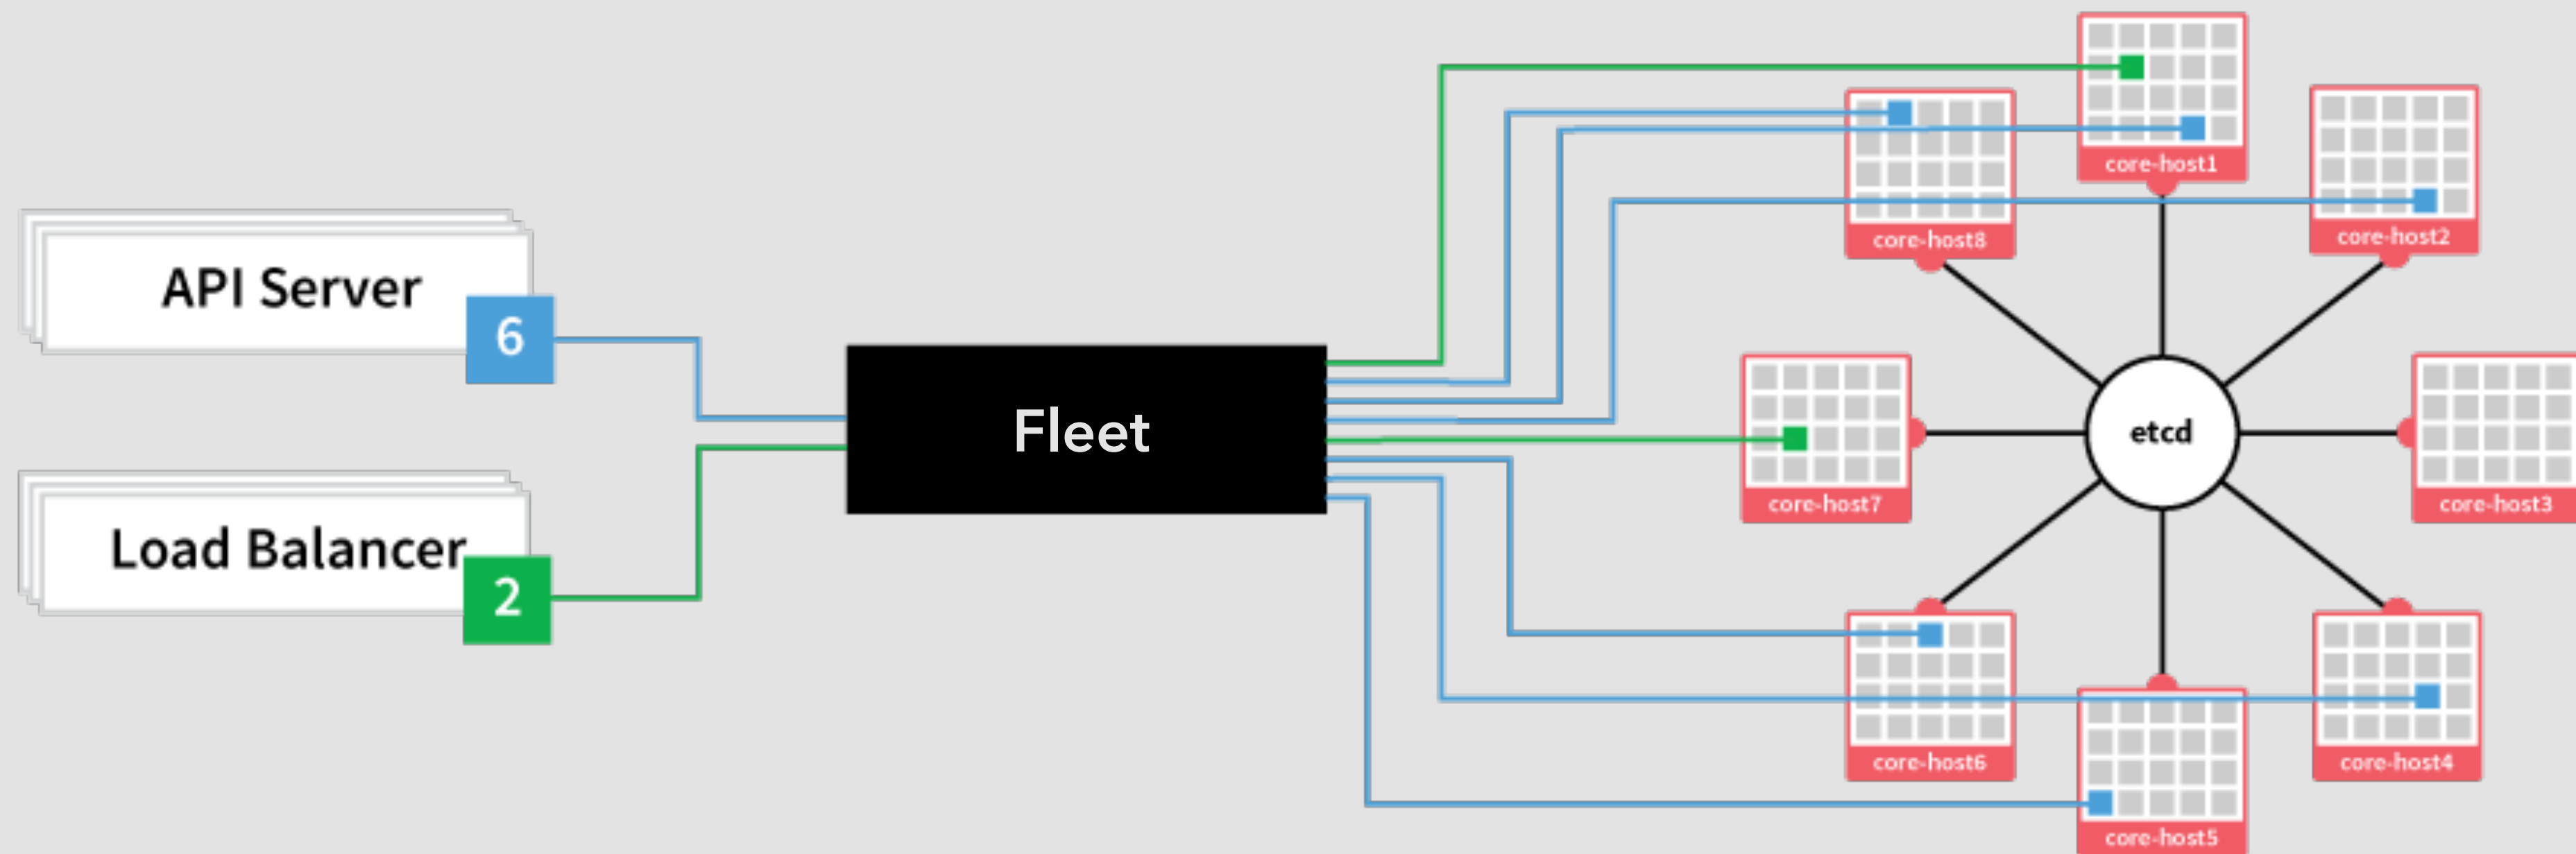
fleet

Etc is a distributed, consistent key-value store for shared configuration and service discovery

- Curl'able user-facing API (HTTP+JSON)
- Optional SSL client cert authentication
- Benchmarked 1000s of writes/s per instance
- Distributed using Raft consensus algorithm

Fleet ties together systemd and **etcd** into a simple distributed init system.

Think of it as systemd that operates at the cluster level instead of the machine level.



Confd - Manage local application configuration files using templates and data from etcd or consul

This

```
upstream sweetapp_pool {  
    {{ range getvs "/services/sweetapp/*" }}  
        server {{ . }};  
    {{ end }}  
}
```

...

```
location /sweetapp/ {  
  
    proxy_pass http://sweetapp/api;
```

...

```
}
```

Becomes This

```
upstream sweetapp_pool {  
    server 10.253.156.164:8811;  
    server 10.253.156.165:8810;  
}
```

...

```
location /sweetapp/ {  
    proxy_pass http://sweetapp/api;
```

...

```
}
```

[Unit]

Description=My Advanced Service

After=etcd.service

After=docker.service

[Service]

TimeoutStartSec=0

ExecStartPre=-/usr/bin/docker kill apache1

ExecStartPre=-/usr/bin/docker rm apache1

ExecStartPre=/usr/bin/docker pull coreos/apache

ExecStart=/usr/bin/docker run --name apache1 -p 8081:80 coreos/apache /usr/sbin

-D FOREGROUND

ExecStartPost=/usr/bin/etcdctl set /domains/example.com/10.10.10.123:8081 run

ExecStop=/usr/bin/docker stop apache1

ExecStopPost=/usr/bin/etcdctl rm /domains/example.com/10.10.10.123:8081

[X-Fleet]

Conflicts=%p@*.service

[Unit]

Description=My Advanced Service

After=etcd.service

After=docker.service

[Service]

TimeoutStartSec=0

ExecStartPre=-/usr/bin/docker kill apache1

ExecStartPre=-/usr/bin/docker rm apache1

ExecStartPre=/usr/bin/docker pull coreos/apache

ExecStart=/usr/bin/docker run --name apache1 -p 8081:80 coreos/apache /usr/sbin

-D FOREGROUND

ExecStartPost=/usr/bin/etcdctl set /domains/example.com/10.10.10.123:8081 run

ExecStop=/usr/bin/docker stop apache1

ExecStopPost=/usr/bin/etcdctl rm /domains/example.com/10.10.10.123:8081

[X-Fleet]

Conflicts=%p@*.service

[Unit]

Description=My Advanced Service

After=etcd.service

After=docker.service

[Service]

TimeoutStartSec=0

ExecStartPre=-/usr/bin/docker kill apache1

ExecStartPre=-/usr/bin/docker rm apache1

ExecStartPre=/usr/bin/docker pull coreos/apache

ExecStart=/usr/bin/docker run --name apache1 -p 8081:80 coreos/apache /usr/sbin

-D FOREGROUND

ExecStartPost=/usr/bin/etcdctl set /domains/example.com/10.10.10.123:8081 run

ExecStop=/usr/bin/docker stop apache1

ExecStopPost=/usr/bin/etcdctl rm /domains/example.com/10.10.10.123:8081

[X-Fleet]

Conflicts=%p@*.service

[Unit]

Description=My Advanced Service

After=etcd.service

After=docker.service

[Service]

TimeoutStartSec=0

ExecStartPre=-/usr/bin/docker kill apache1

ExecStartPre=-/usr/bin/docker rm apache1

ExecStartPre=/usr/bin/docker pull coreos/apache

ExecStart=/usr/bin/docker run --name apache1 -p 8081:80 coreos/apache /usr/sbin/httpd
-D FOREGROUND

ExecStartPost=/usr/bin/etcdctl set /domains/example.com/10.10.10.123:8081 run

ExecStop=/usr/bin/docker stop apache1

ExecStopPost=/usr/bin/etcdctl rm /domains/example.com/10.10.10.123:8081

[X-Fleet]

Conflicts=%p@*.service

[Unit]

Description=My Advanced Service

After=etcd.service

After=docker.service

[Service]

TimeoutStartSec=0

ExecStartPre=-/usr/bin/docker kill apache1

ExecStartPre=-/usr/bin/docker rm apache1

ExecStartPre=/usr/bin/docker pull coreos/apache

ExecStart=/usr/bin/docker run --name apache1 -p 8081:80 coreos/apache /usr/sbin

-D FOREGROUND

ExecStartPost=/usr/bin/etcdctl set /domains/example.com/10.10.10.123:8081 run

ExecStop=/usr/bin/docker stop apache1

ExecStopPost=/usr/bin/etcdctl rm /domains/example.com/10.10.10.123:8081

[X-Fleet]

Conflicts=%p@*.service

[Unit]

Description=My Advanced Service

After=etcd.service

After=docker.service

[Service]

TimeoutStartSec=0

ExecStartPre=-/usr/bin/docker kill apache1

ExecStartPre=-/usr/bin/docker rm apache1

ExecStartPre=/usr/bin/docker pull coreos/apache

ExecStart=/usr/bin/docker run --name apache1 -p 8081:80 coreos/apache /usr/sbin

-D FOREGROUND

ExecStartPost=/usr/bin/etcdctl set /domains/example.com/10.10.10.123:8081 run

ExecStop=/usr/bin/docker stop apache1

ExecStopPost=/usr/bin/etcdctl rm /domains/example.com/10.10.10.123:8081

[X-Fleet]

Conflicts=%p@*.service

[Unit]

Description=My Advanced Service

After=etcd.service

After=docker.service

[Service]

TimeoutStartSec=0

ExecStartPre=-/usr/bin/docker kill apache1

ExecStartPre=-/usr/bin/docker rm apache1

ExecStartPre=/usr/bin/docker pull coreos/apache

ExecStart=/usr/bin/docker run --name apache1 -p 8081:80 coreos/apache /usr/sbin

-D FOREGROUND

ExecStartPost=/usr/bin/etcdctl set /domains/example.com/10.10.10.123:8081 run

ExecStop=/usr/bin/docker stop apache1

ExecStopPost=/usr/bin/etcdctl rm /domains/example.com/10.10.10.123:8081

[X-Fleet]

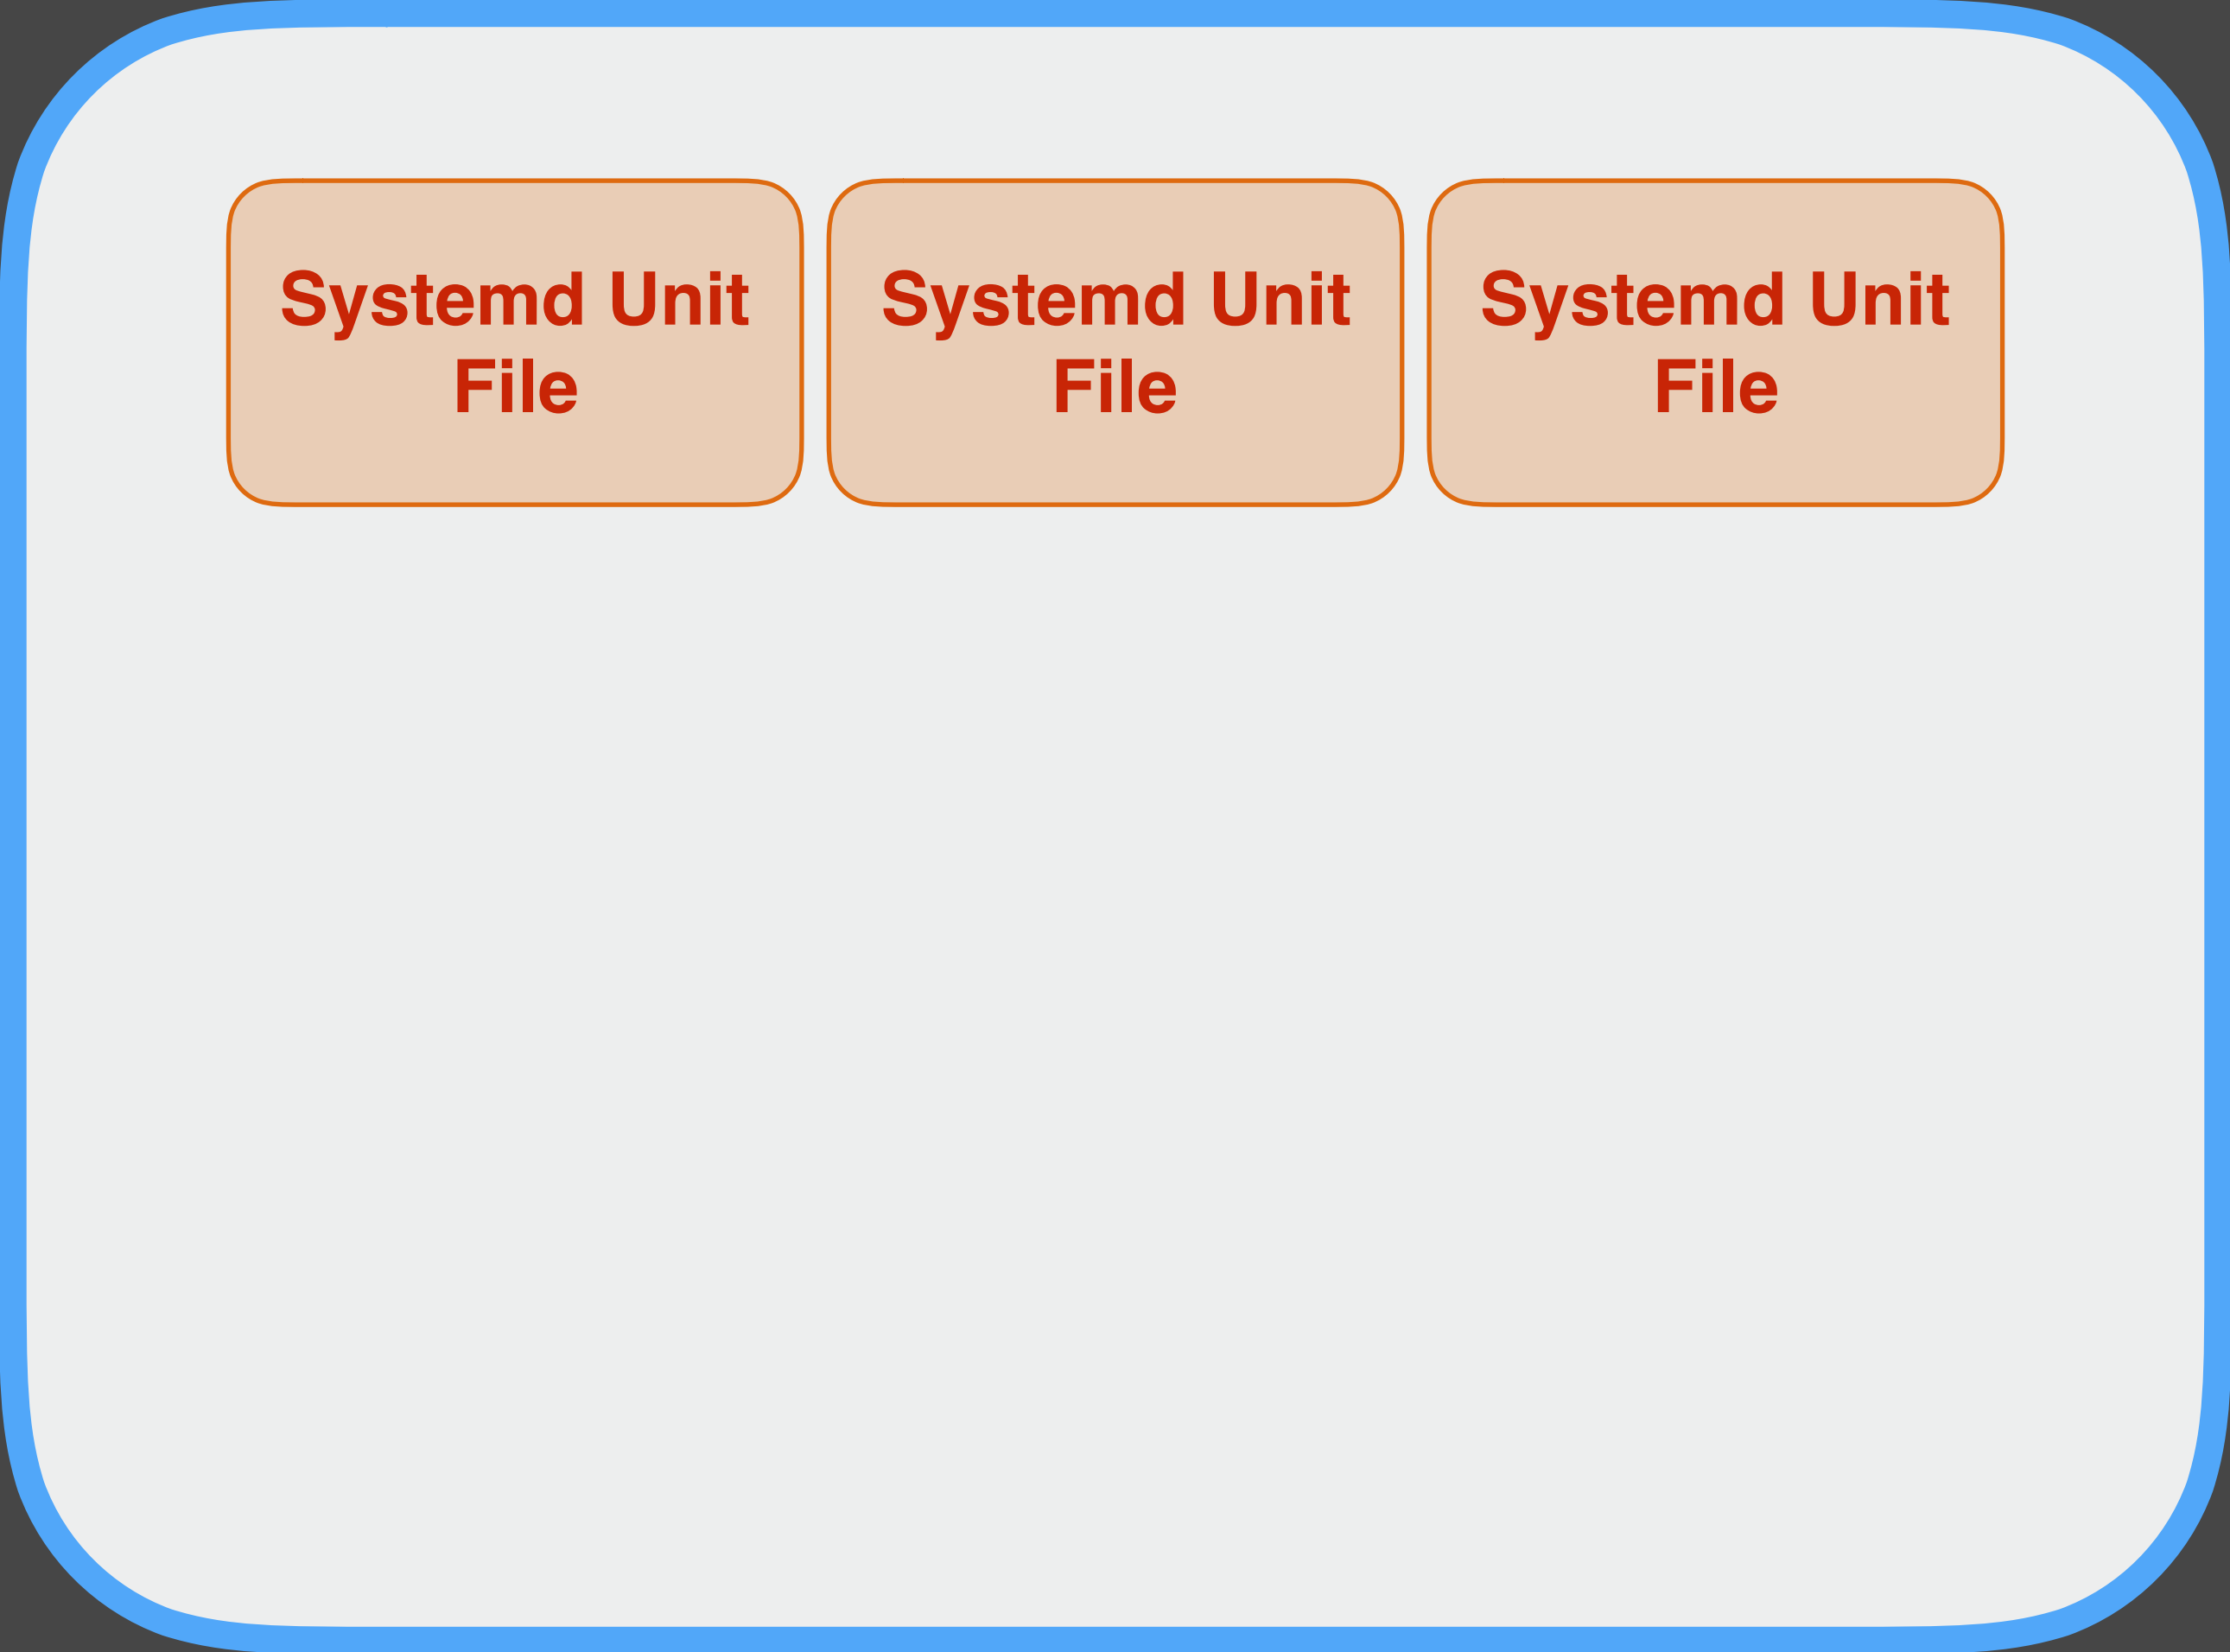
Conflicts=%p@*.service

Option Name	Description
MachineID	Require the unit be scheduled to the machine identified by the given string
MachineOf	Limit eligible machines to the one that hosts a specific unit
MachineMetadata	Limit eligible machines to those with this specific metadata
Conflicts	Prevent a unit from being colocated with other specific units

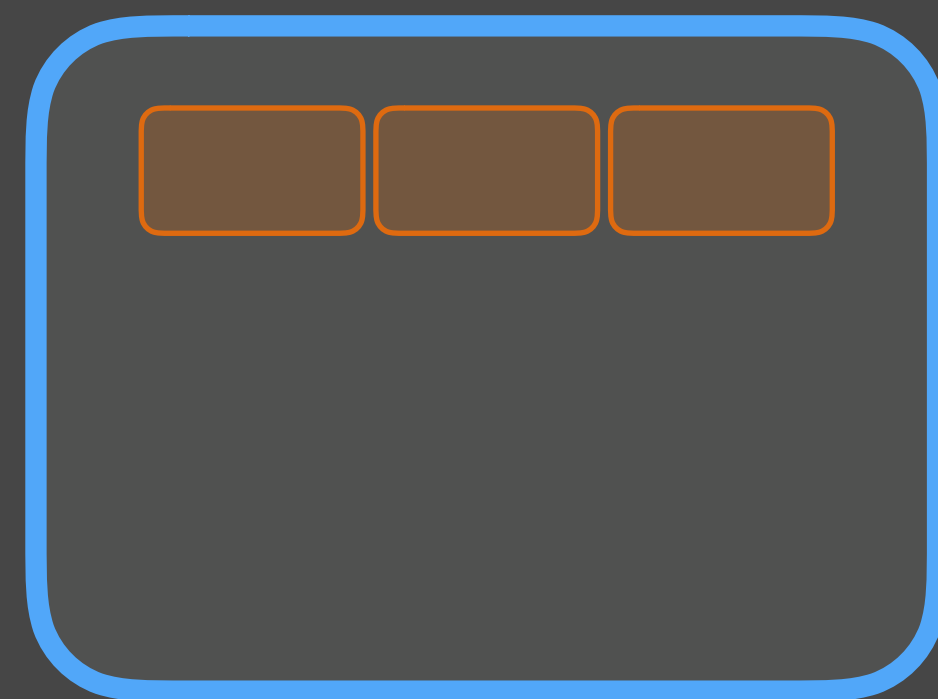
Application Deploys

PUT json version of unit file to **fleet** API


```
curl -s -X PUT -d @unit.json -H 'Content-Type:  
application/json' coreos.demo.jpw.com:49153/fleet/v1/  
units/mycoolapp@12.service
```

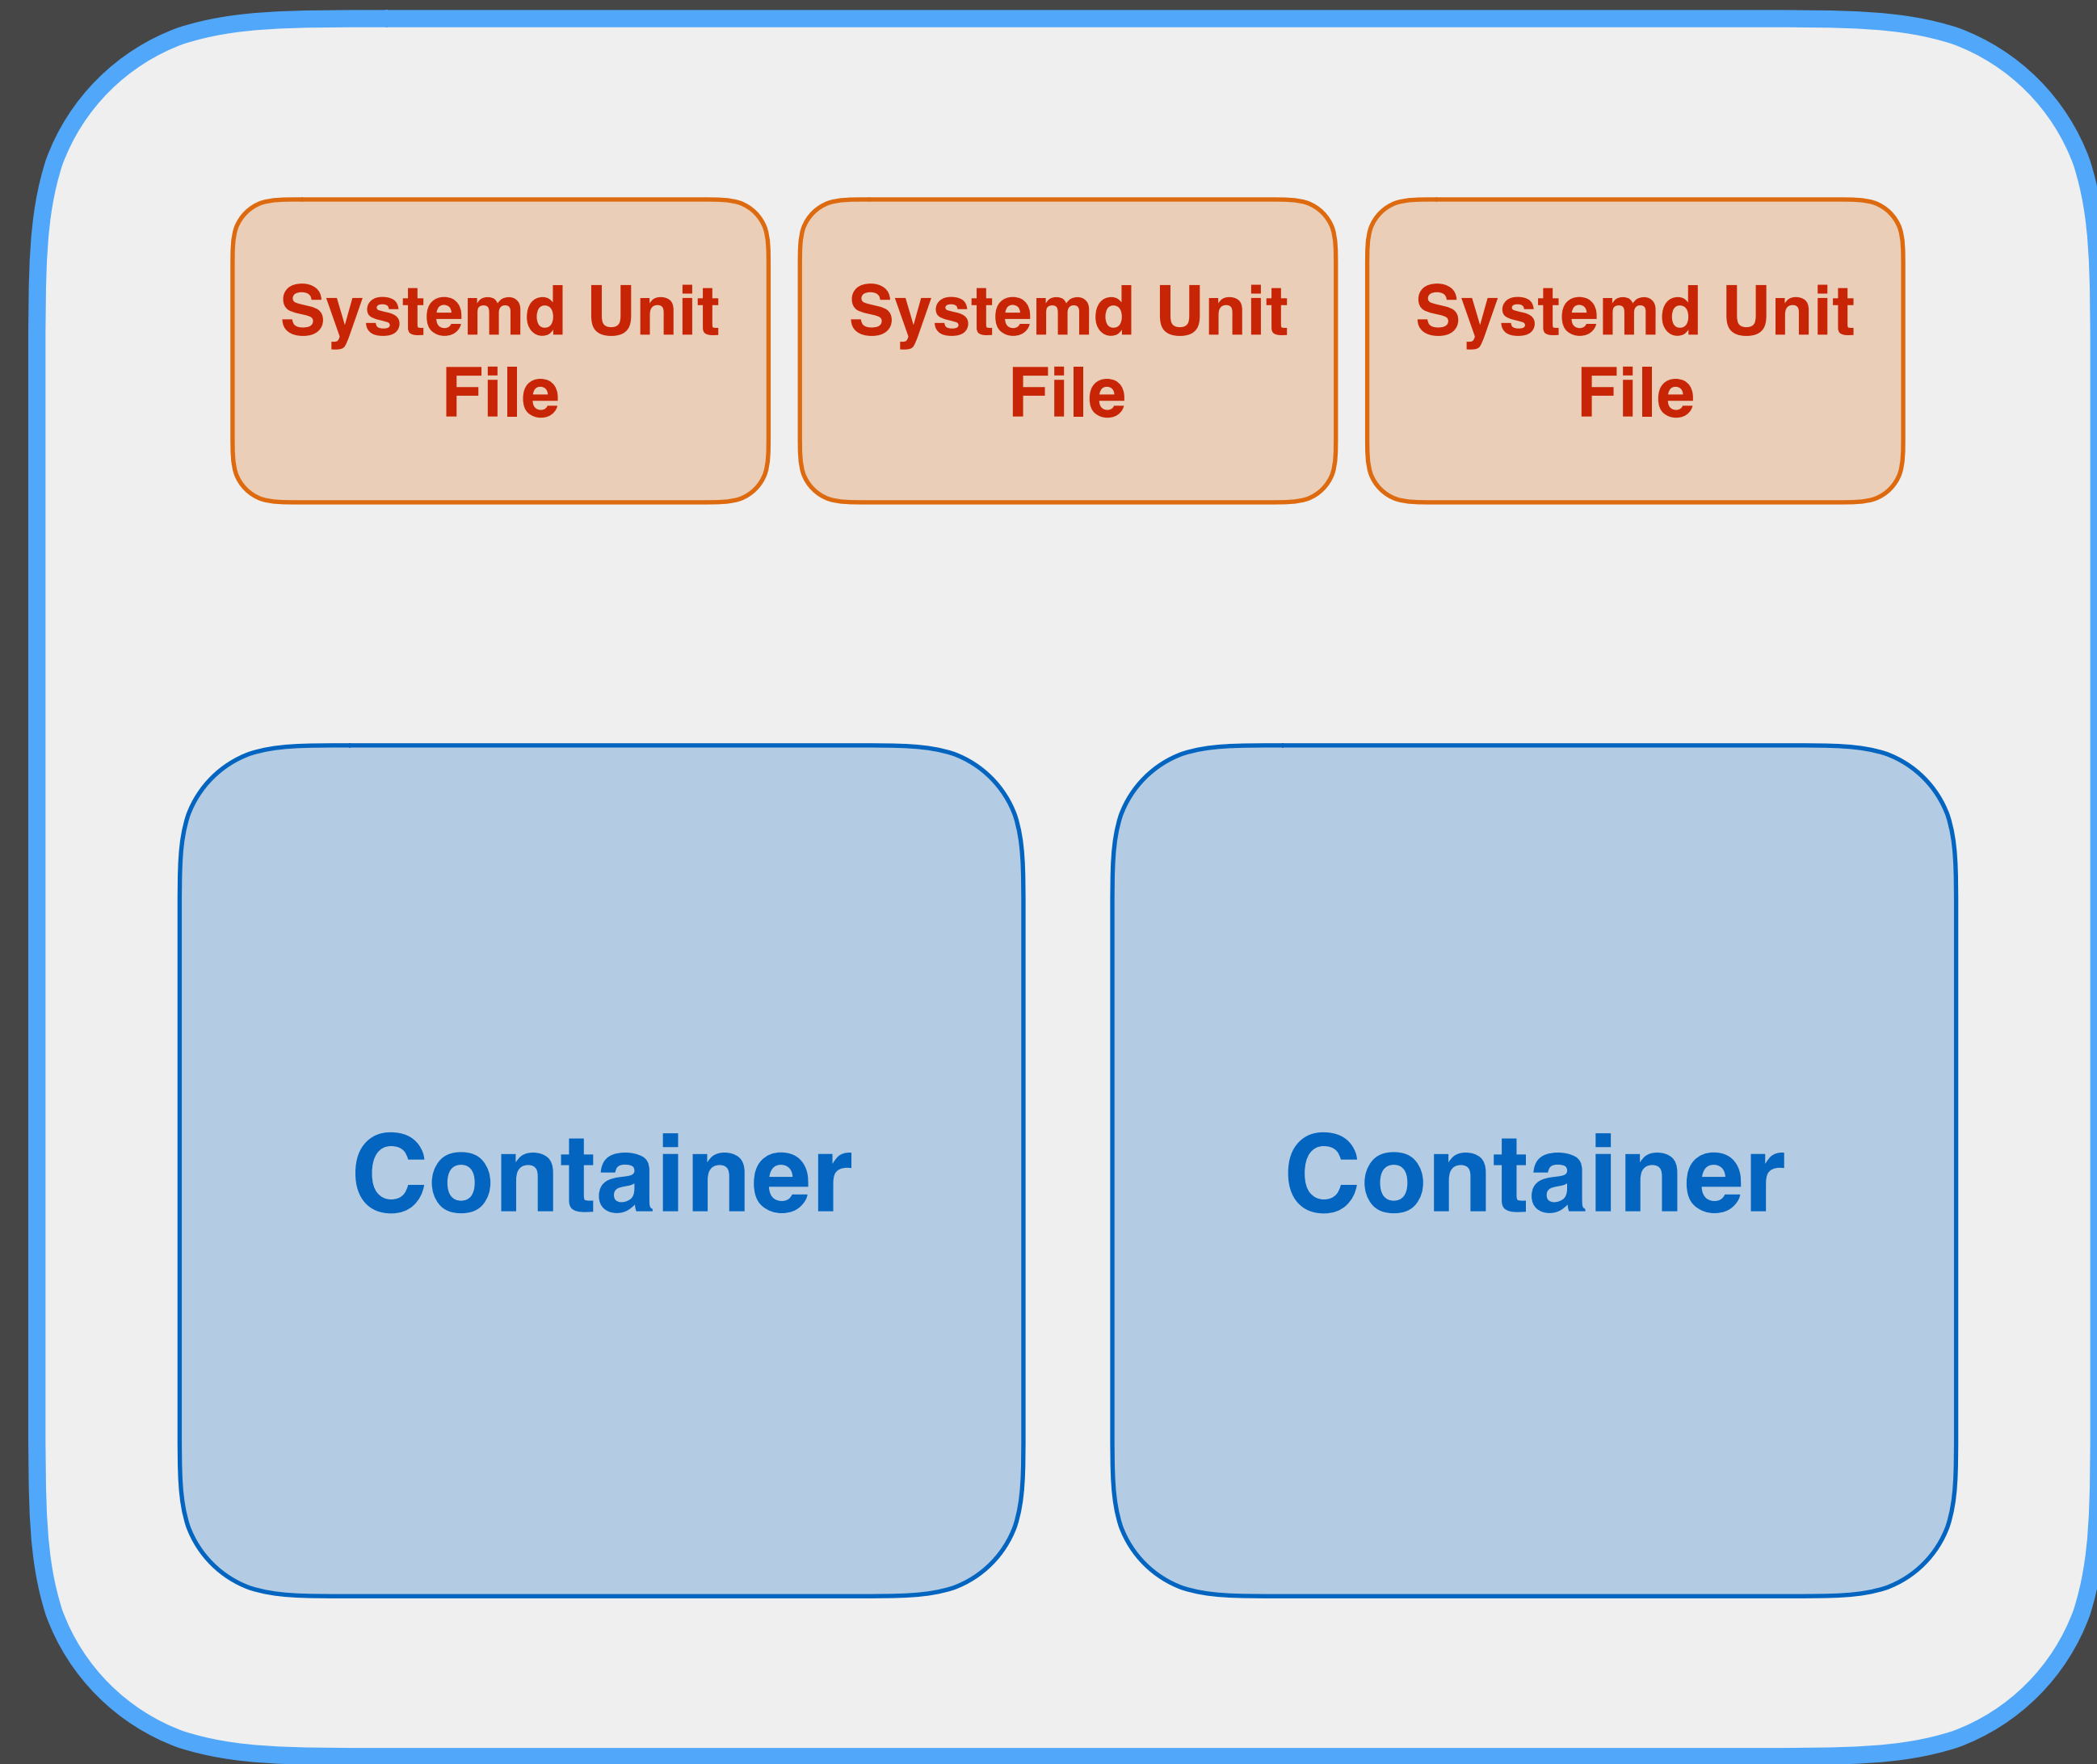


CoreOS Node



CoreOS Node





CoreOS Node

Exposing 10.10.2.3 : 8810

Exposing 10.10.2.3 : 443

Port 8810



Port 8080

172.17.42.2

Port 443



Port 443

172.17.42.3

CoreOS Node | IP 10.10.2.3

Exposing 10.10.2.3 : 8810

Exposing 10.10.2.3 : 443

Port 8810



Port 8080

172.17.42.2

Port 443



Port 443

172.17.42.3

CoreOS Node | IP 10.10.2.3

Exposing 10.10.2.3 : 8810

Exposing 10.10.2.3 : 443

Port 8810



Port 8080

172.17.42.2

Port 443



Port 443

172.17.42.3

CoreOS Node | IP 10.10.2.3

Exposing 10.10.2.3 : 8810

Exposing 10.10.2.3 : 443

Port 8810



Port 8080

172.17.42.2

Port 443



Port 443

172.17.42.3

CoreOS Node | IP 10.10.2.3

Service Discovery and Configuration


```
# etcdctl ls /services/quad  
/services/quad/10.253.156.164  
/services/quad/10.253.156.165
```

```
# etcdctl ls /services/quad  
/services/quad/10.253.156.164  
/services/quad/10.253.156.165
```

```
# etcdctl get /services/quad/10.253.156.165
```

```
# etcdctl ls /services/quad  
/services/quad/10.253.156.164  
/services/quad/10.253.156.165
```

```
# etcdctl get /services/quad/10.253.156.165  
10.253.156.165:8810
```

Hey Jim, what is a sidekick container?

A container that performs other duties that are related to our main container but shouldn't be directly built into that application

Nginx Presence Unit File

[Unit]

BindsTo=nginx-1.0.0@%i.service

[Service]

EnvironmentFile=/etc/environment

ExecStartPre=-/usr/bin/docker kill nginx-presence-%i

ExecStartPre=-/usr/bin/docker rm nginx-presence-%i

ExecStart=/usr/bin/docker run --rm --name nginx-presence-%i -e AWS_ACCESS_KEY=foo -e AWS_SECRET_KEY=bar -e AWS_REGION=us-east-1 -e ELB_NAME=coreos-appname -e PROXY_HOST=10.10.11.11 -e PROXY_PORT=3128 jpweb/nginx-presence

ExecStop=/usr/bin/docker stop nginx-presence-%i

[X-Fleet]

MachineOf=nginx-1.0.0-secure@%i.service

Nginx Presence Unit File

[Unit]

BindsTo=nginx-1.0.0@%i.service

[Service]

EnvironmentFile=/etc/environment

ExecStartPre=-/usr/bin/docker kill nginx-presence-%i

ExecStartPre=-/usr/bin/docker rm nginx-presence-%i

ExecStart=/usr/bin/docker run --rm --name nginx-presence-%i -e AWS_ACCESS_KEY=foo -e AWS_SECRET_KEY=bar -e AWS_REGION=us-east-1 -e ELB_NAME=coreos-appname -e PROXY_HOST=10.10.11.11 -e PROXY_PORT=3128 jpweb/nginx-presence

ExecStop=/usr/bin/docker stop nginx-presence-%i

[X-Fleet]

MachineOf=nginx-1.0.0-secure@%i.service

Nginx Presence Unit File

[Unit]

Bindsto=nginx-1.0.0@%i.service

[Service]

EnvironmentFile=/etc/environment

ExecStartPre=-/usr/bin/docker kill nginx-presence-%i

ExecStartPre=-/usr/bin/docker rm nginx-presence-%i

ExecStart=/usr/bin/docker run --rm --name nginx-presence-%i -e AWS_ACCESS_KEY=foo -e AWS_SECRET_KEY=bar -e AWS_REGION=us-east-1 -e ELB_NAME=coreos-appname -e PROXY_HOST=10.10.11.11 -e PROXY_PORT=3128 jpweb/nginx-presence

ExecStop=/usr/bin/docker stop nginx-presence-%i

[X-Fleet]

MachineOf=nginx-1.0.0-secure@%i.service

SDN

Software-defined Networking



- Containers IP address allocation across cluster
- Dynamic dns across cluster.
- Encrypted connections
- Easy install short up and running time
- Container networking that runs in a container?



**YO DAWG! I HEARD YOU LIKE
NETWORKS**

**SO I PUT AN SDN IN A CONTAINER, IN A VM, THAT IS A PART OF AN SDN, ON A HOST, ON A
NETWORK, ON THE INTERNET**

memegenerator.net

This time, yeah really

(on aws anyway)

Weave

ethwe: 192.168.255.0/22

Container

eth0: 172.17.0.2/16

Docker Environment

docker0: 172.17.0.2/16

Running Instance

eth0: 10.10.1.1/12

AWS

another SDN, etc

Lessons Learned

Reference Links

- <https://coreos.com>
- <https://www.docker.com>
- <https://coreos.com/etcd/docs/latest/>
- <https://github.com/kelseyhightower/confd>
- <https://raft.github.io>
- @jpw
- <https://github.com/jpweber>